Mirror Descent

Axel Böhm

November 7, 2021

Recap on (sub)-gradient descent

◇ When we used a norm $\| \cdot \|$ we meant the 2-*norm*, i.e.

$$\|x\|_2 = \sqrt{x_1^2 + x_2^2 + \cdots + x_d^2}.$$

◇ In **gradient descent** we used Lipschitz continuity:

$$\|\nabla f(x) - \nabla f(y)\| \le L\|x - y\|.$$

(Lead to a complexity of $\mathcal{O}(\frac{L}{k})$)

◇ **Sub-gradient descent**: used $\|g\| \le G$ which lead to $\mathcal{O}(\frac{G}{\sqrt{k}})$.

◇ But there are other norms $\|x\|_\infty \le \|x\|_2 \le \sqrt{d}\|x\|_\infty$.
It can happen that $\|g\|_\infty \le G$ but $\|g\|_2 \approx \sqrt{d}\,G$.

we lose dimension independence

Recap on (sub)-gradient descent

◇ When we used a norm $\| \cdot \|$ we meant the 2-*norm*, i.e.

$$\|x\|_2 = \sqrt{x_1^2 + x_2^2 + \cdots + x_d^2}.$$

◇ In **gradient descent** we used Lipschitz continuity:

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|.$$

(Lead to a complexity of $\mathcal{O}(\frac{L}{k})$)

◇ **Sub-gradient descent**: used $\|g\| \leq G$ which lead to $\mathcal{O}(\frac{G}{\sqrt{k}})$.

◇ But there are other norms $\|x\|_\infty \leq \|x\|_2 \leq \sqrt{d}\|x\|_\infty$.
It can happen that $\|g\|_\infty \leq G$ but $\|g\|_2 \approx \sqrt{d}G$.

> we lose dimension independence

Different norms?

*But where did we use the norm in the **method**?*

Gradient Descent

$$x_{k+1} = x_k - \alpha \nabla f(x_k)$$

equivalently

$$x_{k+1} = \underset{x \in \mathbb{R}^d}{\arg \min} \left\{ f(x_k) + \langle \nabla f(x_k), x - x_k \rangle + \frac{1}{2\alpha} \|x - x_k\|_2^2 \right\}$$

We can replace the 2-norm with a more general **distance**.

Bregman distance

$h : \mathbb{R}^d \to \mathbb{R} \cup \{+\infty\}$ is convex

(i) h is differentiable of the interior of dom h

(ii) h is 1-strongly convex w.r.t. $\| \cdot \|_2$

Then
$$\mathcal{D}_h(x, y) := h(x) - h(y) - \langle \nabla h(y), x - y \rangle.$$

Properties

$\diamond \ \mathcal{D}_h(x, y) \geq 0$

$\diamond \ \mathcal{D}_h(x, y) \neq \mathcal{D}_h(y, x)$

$\diamond \ \mathcal{D}_h(\cdot, y)$ is convex for all y

$$\mathcal{D}_h(x, y) \approx \frac{1}{2} \langle \nabla^2 h(y)(x - y), x - y \rangle = \frac{1}{2} \|x - y\|_{\nabla^2 h(y)}^2$$

$\diamond \ \mathcal{D}_h(x, y) \geq \frac{1}{2} \|x - y\|^2$ (1-strong convexity)

Bregman distance

$h : \mathbb{R}^d \to \mathbb{R} \cup \{+\infty\}$ is convex
  (i) h is differentiable of the interior of dom h
 (ii) h is 1-strongly convex w.r.t. $\| \cdot \|_2$
Then
$$\mathcal{D}_h(x, y) := h(x) - h(y) - \langle \nabla h(y), x - y \rangle.$$

Properties

$\diamond$ $\mathcal{D}_h(x, y) \geq 0$

$\diamond$ $\mathcal{D}_h(x, y) \neq \mathcal{D}_h(y, x)$

$\diamond$ $\mathcal{D}_h(\cdot, y)$ is convex for all y

$$\mathcal{D}_h(x, y) \approx \frac{1}{2} \langle \nabla^2 h(y)(x - y), x - y \rangle = \frac{1}{2} \|x - y\|_{\nabla^2 h(y)}^2$$

$\diamond$ $\mathcal{D}_h(x, y) \geq \frac{1}{2} \|x - y\|^2$ (1-strong convexity)

Examples

◇ $h(x) = \frac{1}{2}\|x\|_2^2$ gives $\mathcal{D}_h(x, y) = \|x - y\|^2$

◇ $h(x) = \frac{1}{2(p-1)}\|x\|_p^2$ with $p \in [1, 2]$

◇ $\Delta^d = \{x \in \mathbb{R}_+^d : \sum_{i=1}^d x_i = 1\}$ the *unit simplex* and

$$h(x) = \begin{cases} \sum_{i=1}^d x_i \log(x_i) & x_i > 0 \\ +\infty & \text{otherwise} \end{cases}$$

the **Negative entropy**.

Negative entropy

◇ Negative entropy: $h(x) = \sum_{i=1}^{d} x_i \log(x_i)$ for $x_i > 0$.

◇ Then $\nabla h(x) = \log(x) + 1$ (coordinatewise) and

$$\mathcal{D}_h(x, y) = \sum_{i=1}^{d} x_i \log(x_i) - y_i \log(y_i) - \langle \log(y) + 1, x - y \rangle$$

$$= \sum_{i=1}^{d} x_i \log(x_i) - \sum_{i=1}^{d} x_i \log(y_i)$$

$$= \sum_{i=1}^{d} x_i \log\left(\frac{x_i}{y_i}\right)$$

Known as Kullback-Leibler divergence $K(X \| Y)$.

◇ Is strongly convex over Δ

$$\mathcal{D}(x, y) \geq \frac{1}{2} \|x - y\|_1^2 \quad \text{Pinsker's ineq.}$$

Mirror descent

Idea: replace squared Euclidian norm with more general object:

$$x_{k+1} = \underset{x\in\mathbb{R}^d}{\arg\min}\{f(x_k) + \langle\nabla f(x_k), x - x_k\rangle + \frac{1}{\alpha}\mathcal{D}_h(x, x_k)\}$$

$$= \underset{x\in\mathbb{R}^d}{\arg\min}\{\langle\nabla f(x_k), x\rangle + \frac{1}{\alpha}\mathcal{D}_h(x, x_k)\}$$

$$= \underset{x\in\mathbb{R}^d}{\arg\min}\{\langle\nabla f(x_k), x\rangle + \frac{1}{\alpha}(h(x) - h(x_k) - \langle\nabla h(x_k), x - x_k\rangle)\}$$

$$= \underset{x\in\mathbb{R}^d}{\arg\min}\{\langle\nabla f(x_k), x\rangle + \frac{1}{\alpha}(h(x) - \langle\nabla h(x_k), x\rangle)\}$$

$$= \underset{x\in\mathbb{R}^d}{\arg\min}\{\langle\alpha\nabla f(x_k) - \nabla h(x_k), x\rangle + h(x)\}$$

Question: But why *mirror* descent?

Mirror descent

Idea: replace squared Euclidian norm with more general object:

$$
\begin{aligned}
x_{k+1} &= \arg\min_{x\in\mathbb{R}^d}\{f(x_k) + \langle\nabla f(x_k), x - x_k\rangle + \frac{1}{\alpha}\mathcal{D}_h(x, x_k)\} \\
&= \arg\min_{x\in\mathbb{R}^d}\{\langle\nabla f(x_k), x\rangle + \frac{1}{\alpha}\mathcal{D}_h(x, x_k)\} \\
&= \arg\min_{x\in\mathbb{R}^d}\{\langle\nabla f(x_k), x\rangle + \frac{1}{\alpha}(h(x) - h(x_k) - \langle\nabla h(x_k), x - x_k\rangle)\} \\
&= \arg\min_{x\in\mathbb{R}^d}\{\langle\nabla f(x_k), x\rangle + \frac{1}{\alpha}(h(x) - \langle\nabla h(x_k), x\rangle)\} \\
&= \arg\min_{x\in\mathbb{R}^d}\{\langle\alpha\nabla f(x_k) - \nabla h(x_k), x\rangle + h(x)\}
\end{aligned}
$$

Question: But why *mirror* descent?

Mirror descent

Idea: replace squared Euclidian norm with more general object:

$$
\begin{aligned}
x_{k+1} &= \underset{x \in \mathbb{R}^d}{\arg\min}\{f(x_k) + \langle \nabla f(x_k), x - x_k \rangle + \frac{1}{\alpha}\mathcal{D}_h(x, x_k)\} \\
&= \underset{x \in \mathbb{R}^d}{\arg\min}\{\langle \nabla f(x_k), x \rangle + \frac{1}{\alpha}\mathcal{D}_h(x, x_k)\} \\
&= \underset{x \in \mathbb{R}^d}{\arg\min}\{\langle \nabla f(x_k), x \rangle + \frac{1}{\alpha}(h(x) - h(x_k) - \langle \nabla h(x_k), x - x_k \rangle)\} \\
&= \underset{x \in \mathbb{R}^d}{\arg\min}\{\langle \nabla f(x_k), x \rangle + \frac{1}{\alpha}(h(x) - \langle \nabla h(x_k), x \rangle)\} \\
&= \underset{x \in \mathbb{R}^d}{\arg\min}\{\langle \alpha \nabla f(x_k) - \nabla h(x_k), x \rangle + h(x)\}
\end{aligned}
$$

Question: But why *mirror* descent?

The Mirror part

$$x_{k+1} = \underset{x \in \mathbb{R}^d}{\arg\min} \{\langle \alpha \nabla f(x_k) - \nabla h(x_k), x \rangle + h(x)\}$$

By optimality condition:

$$0 = \alpha \nabla f(x_k) - \nabla h(x_k) + \nabla h(x_{k+1})$$

Therefore

$$\nabla h(x_{k+1}) = \nabla h(x_k) - \alpha \nabla f(x_k)$$
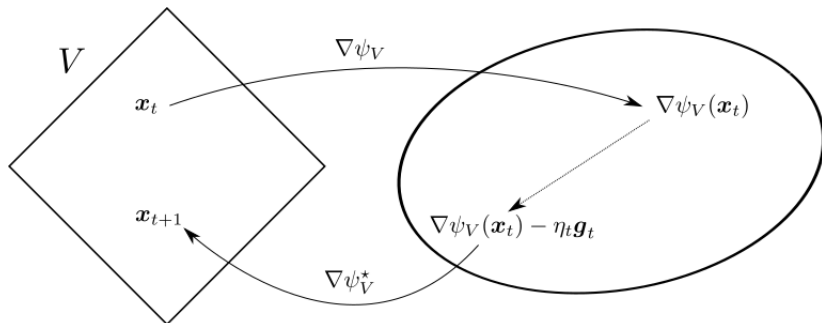
Why it's called mirror descent

Figure: $\psi = h$

Mirror Descent w.r.t. neg. entropy

Negative entropy: $h(x) = \sum_{i=1}^{d} x_i \log(x_i)$ for $x_i > 0$.
We define $a := \alpha \nabla f(x_k) - \nabla h(x_k)$. Then

$$x_{k+1} = \underset{x \in \Delta}{\arg\min}\{\langle a, x \rangle + h(x)\}$$

Results in

$$x_{k+1} = x_k \cdot * \, e^{-\alpha \nabla f(x_k)}$$

in a pointwise sense.

(General) mirror descent convergence statement

Since we changed norm in the space of the variable x, we need to go to the dual norms in the space of the subgradients

$$\|y\|_* := \max_{\|x\|=1} \{\langle y, x \rangle\}.$$

Theorem

In $(\mathbb{R}^d, \|\cdot\|)$ and subgradients bounded in dual norm $\|g_k\|_* \leq G$, then

$$f(\bar{x}_k) - f^* \leq \frac{(\mathcal{D}(x^*, x_0))^{1/2} G}{\sqrt{k}},$$

where $\bar{x}_k$ denotes the averaged iterates, as usual.

Convergence w.r.t. neg. entropy

What about $\mathcal{D}(x^*, x_0)$? Let $x_0 = (\frac{1}{n}, \cdots, \frac{1}{n})$, then

$$\mathcal{D}(x, x_0) = \sum x_i \log\left(\frac{x_i}{\frac{1}{n}}\right) = \sum x_i \log(x_i) + \log(n) \le \log(n)$$

while $\|x_0 - x^*\|^2 \le 2$.
But if

$$\|g\|_\infty = \|g\|_1^* \le G$$

we can still have

$$\|g\|_2 \approx \sqrt{d}\, G.$$

Proof

In the Euclidian space we used

$$\langle x_{k+1} - x_k, x^* - x_{k+1} \rangle$$
$$= \frac{1}{2} \|x^* - x_k\|^2 - \frac{1}{2} \|x^* - x_{k+1}\|^2 - \frac{1}{2} \|x_{k+1} - x_k\|^2.$$

Similar 3-point identity holds for Bregman distances:

$$\langle \nabla h(x_{k+1}) - \nabla h(x_k), x^* - x_{k+1} \rangle =$$
$$= D(x^*, x_k) - D(x^*, x_{k+1}) - D(x_{k+1}, x_k).$$

Therefore

$$D(x^*, x_{k+1}) \leq D(x^*, x_k) - D(x_{k+1}, x_k) + \alpha \langle g_k, x^* - x_{k+1} \rangle.$$

Proof II

$$D(x^*, x_{k+1}) \leq D(x^*, x_k) - D(x_{k+1}, x_k) + \alpha \langle g_k, x^* - x_{k+1} \rangle.$$

Last term is not quite right.

$$\begin{aligned} \langle g_k, x^* - x_{k+1} \rangle &= \langle g_k, x^* - x_k \rangle + \langle g_k, x_k - x_{k+1} \rangle \\ &\leq f(x^*) - f(x_k) + \|g_k\|_* \|x_k - x_{k+1}\| \\ &\leq f(x^*) - f(x_k) + \frac{\alpha \|g_k\|_*^2}{2} + \frac{\|x_k - x_{k+1}\|^2}{2\alpha}. \end{aligned}$$

Combined we get that

$$\begin{aligned} D(x^*, x_{k+1}) \leq\ &D(x^*, x_k) - D(x_{k+1}, x_k) + \alpha(f(x^*) - f(x_k)) \\ &+ \frac{\alpha^2 \|g_k\|_*^2}{2} + \frac{\|x_k - x_{k+1}\|^2}{2}. \end{aligned}$$

Proof III

We assumed strong convexity of h:

$$D(x_{k+1}, x_k) \geq \frac{1}{2}\|x_{k+1} - x_k\|^2.$$

Yields

$$D(x^*, x_{k+1}) \leq D(x^*, x_k) + \alpha(f(x^*) - f(x_k)) + \frac{\alpha^2 \|g_k\|_*^2}{2}$$

Continue as always

$$\frac{1}{k}\sum_{i=1}^{k} f(x_i) - f^* \leq \frac{D(x^*, x_0)}{\alpha k} \frac{\alpha G^2}{2}$$

What about the smooth case

⬦ Talked about how to get better constants in the "bounded subgradients" setting

⬦ but can't make them bounded if they are not

However,

⬦ Can also come up with a new notion of smoothness

$$f(y) \leq f(x) + \langle \nabla f(x), y - x \rangle + LD(y, x)$$

⬦ which might hold even if f is not smooth in classical sense