

Matrix games

Axel Böhm

October 3, 2021

1 Introduction

2 Algorithms

Introduction

Given

- ◇ Player I (rows, Alice)
- ◇ Player II (columns, Bob)
- ◇ a *payoff* matrix $A \in \mathbb{R}^{m \times n}$

Every round

- (i) Alice picks (row) strategy $i \in [m] := \{1, \dots, m\}$ Bob picks (col) strategy $j \in [n]$
- (ii) Bob pays Alice the amount $a_{i,j}$

zero-sum game

Example: penalty game

A handwritten payoff matrix for a penalty game. The matrix is written on a light yellow grid background. The columns are labeled 'kicker' with 'L' and 'R' below them. The rows are labeled 'goalkeeper' with 'L' and 'R' to the left of them. The payoffs are written in the cells, separated by a vertical line from the row label and a horizontal line from the column label. The payoffs are: (L, L) = 1, (L, R) = -1, (R, L) = -1, (R, R) = 1.

		kicker	
		L	R
goalkeeper	L	1	-1
	R	-1	1

Figure: penalty game

Example: prisoners dilemma

	Confess A	Stay quiet A
Confess B	6 6	10 0
Stay quiet B	0 10	2 2

Figure: prisoners dilemma

Worst case

- ◇ if Alice chooses strategy i she gets (at least): $\min_{j \in [n]} a_{i,j}$
- ◇ Alice can ensure payoff $\max_{i \in [m]} \min_{j \in [n]} a_{i,j}$
- ◇ Bob pays (at most) $\min_{j \in [n]} \max_{i \in [m]} a_{i,j}$

We claim:

$$\max_i \min_j a_{i,j} \leq \min_j \max_i a_{i,j}$$

"Tallest dwarf is not as tall as the smallest giant."

But: No equality in general!

Worst case

- ◇ if Alice chooses strategy i she gets (at least): $\min_{j \in [n]} a_{i,j}$
- ◇ Alice can ensure payoff $\max_{i \in [m]} \min_{j \in [n]} a_{i,j}$
- ◇ Bob pays (at most) $\min_{j \in [n]} \max_{i \in [m]} a_{i,j}$

We claim:

$$\max_i \min_j a_{i,j} \leq \min_j \max_i a_{i,j}$$

"Tallest dwarf is not as tall as the smallest giant."

But: No equality in general!

Worst case

- ◇ if Alice chooses strategy i she gets (at least): $\min_{j \in [n]} a_{i,j}$
- ◇ Alice can ensure payoff $\max_{i \in [m]} \min_{j \in [n]} a_{i,j}$
- ◇ Bob pays (at most) $\min_{j \in [n]} \max_{i \in [m]} a_{i,j}$

We claim:

$$\max_i \min_j a_{i,j} \leq \min_j \max_i a_{i,j}$$

“Tallest dwarf is not as tall as the smallest giant.”

But: No equality in general!

Worst case

- ◇ if Alice chooses strategy i she gets (at least): $\min_{j \in [n]} a_{i,j}$
- ◇ Alice can ensure payoff $\max_{i \in [m]} \min_{j \in [n]} a_{i,j}$
- ◇ Bob pays (at most) $\min_{j \in [n]} \max_{i \in [m]} a_{i,j}$

We claim:

$$\max_i \min_j a_{i,j} \leq \min_j \max_i a_{i,j}$$

“Tallest dwarf is not as tall as the smallest giant.”

But: **No equality in general!**

Proof of the min-max theorem

$$\begin{aligned}a_{ij} &\leq a_{ij} && \forall i, j \\a_{ij} &\leq \max_i a_{ij} && \forall i, j \\ \min_j a_{ij} &\leq \min_j \max_i a_{ij} && \forall i\end{aligned}$$

Definition

We call (i^*, j^*) a saddle point (or *Nash equilibrium*) if

$$\max_i a_{ij^*} = a_{i^*j^*} = \min_j a_{i^*j}.$$

These are called *pure strategies*.

Proof of the min-max theorem

$$a_{ij} \leq a_{ij} \quad \forall i, j$$

$$a_{ij} \leq \max_i a_{ij} \quad \forall i, j$$

$$\min_j a_{ij} \leq \min_j \max_i a_{ij} \quad \forall i$$

Definition

We call (i^*, j^*) a saddle point (or *Nash equilibrium*) if

$$\max_i a_{ij^*} = a_{i^*j^*} = \min_j a_{i^*j}.$$

These are called *pure strategies*.

Proof of the min-max theorem

$$\begin{aligned}a_{ij} &\leq a_{ij} && \forall i, j \\a_{ij} &\leq \max_i a_{ij} && \forall i, j \\ \min_j a_{ij} &\leq \min_j \max_i a_{ij} && \forall i\end{aligned}$$

Definition

We call (i^*, j^*) a saddle point (or *Nash equilibrium*) if

$$\max_i a_{ij^*} = a_{i^*j^*} = \min_j a_{i^*j}.$$

These are called *pure strategies*.

Proof of the min-max theorem

$$\begin{aligned}a_{ij} &\leq a_{ij} && \forall i, j \\a_{ij} &\leq \max_i a_{ij} && \forall i, j \\ \min_j a_{ij} &\leq \min_j \max_i a_{ij} && \forall i\end{aligned}$$

Definition

We call (i^*, j^*) a saddle point (or *Nash equilibrium*) if

$$\max_i a_{ij^*} = a_{i^*j^*} = \min_j a_{i^*j}.$$

These are called *pure strategies*.

Rock paper scissors

	R	P	S
R	0	-1	1
P	1	0	-1
S	-1	1	0

Mixed Strategies

With pure strategies we do not always have a saddle point.

von Neumann (1928) — Mixed strategies

- ◇ Alice picks strategies $1, \dots, m$ with probabilities $x \in \Delta_m$
- ◇ Bob picks strategies $1, \dots, n$ with probabilities $y \in \Delta_n$

Expected gain of Alice is

$$\langle x, Ay \rangle = \sum_{i,j} a_{ij} x_i y_j$$

Theorem

Saddle point exists Expected gain of Alice = expected loss of Bob

$$\max_{x \in \Delta} \min_{y \in \Delta} \langle x, Ay \rangle = \min_{y \in \Delta} \max_{x \in \Delta} \langle x, Ay \rangle.$$

Mixed Strategies

With pure strategies we do not always have a saddle point.

von Neumann (1928) — Mixed strategies

- ◇ Alice picks strategies $1, \dots, m$ with probabilities $x \in \Delta_m$
- ◇ Bob picks strategies $1, \dots, n$ with probabilities $y \in \Delta_n$

Expected gain of Alice is

$$\langle x, Ay \rangle = \sum_{i,j} a_{ij} x_i y_j$$

Theorem

Saddle point exists Expected gain of Alice = expected loss of Bob

$$\max_{x \in \Delta} \min_{y \in \Delta} \langle x, Ay \rangle = \min_{y \in \Delta} \max_{x \in \Delta} \langle x, Ay \rangle.$$

Stopping criteria

$$\min_{x \in \Delta} \max_{y \in \Delta} \langle x, Ay \rangle =: v$$

$$\max_{y \in \Delta} \min_{x \in \Delta} \langle x, Ay \rangle = v$$

stopping criterion

$$f_p(x) - f_p(x^*) = f_p(x) - v \leq \epsilon/2$$

$$f_d(y^*) - f_d(y) = v - f_d(y) \leq \epsilon/2$$

$$\Rightarrow f_p(x) - f_d(y) \leq \epsilon$$

Before we never had the optimal value!

Solution will always be on the boundary:

$$f_p(x) = \max_{y \in \Delta} \langle A^T x, y \rangle = \max_j \langle A^T x, e_j \rangle$$

Stopping criteria

$$\min_{x \in \Delta} \max_{y \in \Delta} \langle x, Ay \rangle =: v$$

$$\max_{y \in \Delta} \min_{x \in \Delta} \langle x, Ay \rangle = v$$

stopping criterion

$$f_p(x) - f_p(x^*) = f_p(x) - v \leq \epsilon/2$$

$$f_d(y^*) - f_d(y) = v - f_d(y) \leq \epsilon/2$$

$$\Rightarrow f_p(x) - f_d(y) \leq \epsilon$$

Before we never had the optimal value!

Solution will always be on the boundary:

$$f_p(x) = \max_{y \in \Delta} \langle A^T x, y \rangle = \max_j \langle A^T x, e_j \rangle$$

Stopping criteria

$$\min_{x \in \Delta} \max_{y \in \Delta} \langle x, Ay \rangle =: v$$

$$\max_{y \in \Delta} \min_{x \in \Delta} \langle x, Ay \rangle = v$$

stopping criterion

$$f_p(x) - f_p(x^*) = f_p(x) - v \leq \epsilon/2$$

$$f_d(y^*) - f_d(y) = v - f_d(y) \leq \epsilon/2$$

$$\Rightarrow f_p(x) - f_d(y) \leq \epsilon$$

Before we never had the optimal value!

Solution will always be on the boundary:

$$f_p(x) = \max_{y \in \Delta} \langle A^T x, y \rangle = \max_j \langle A^T x, e_j \rangle$$

Consider

$$\min_{x \in \Delta} \max_{y \in \Delta} \langle x, Ay \rangle$$

as a minimization problem

$$\min_{x \in \Delta} f_p(x) = \langle x, Ay^* \rangle.$$

Then, by the **first-order optimality condition**

$$x^* \in \arg \min_{x \in \Delta} f_p(x) \Leftrightarrow \langle \nabla f_p(x^*), x - x^* \rangle \geq 0 \quad \forall x \in \Delta$$

Thus

$$\langle A^T y^*, x - x^* \rangle \geq 0 \quad \forall x \in \Delta$$

$$\langle -Ax^*, y - y^* \rangle \geq 0 \quad \forall y \in \Delta$$

Concatenate the two conditions to get

$$\left\langle \begin{bmatrix} 0 & A^T \\ -A & 0 \end{bmatrix} \begin{pmatrix} x^* \\ y^* \end{pmatrix}, \begin{pmatrix} x \\ y \end{pmatrix} - \begin{pmatrix} x^* \\ y^* \end{pmatrix} \right\rangle \geq 0.$$

Games as Variational Inequalities

We had:

$$\left\langle \begin{bmatrix} 0 & A^T \\ -A & 0 \end{bmatrix} \begin{pmatrix} x^* \\ y^* \end{pmatrix}, \begin{pmatrix} x \\ y \end{pmatrix} - \begin{pmatrix} x^* \\ y^* \end{pmatrix} \right\rangle \geq 0.$$

By rewriting $z = (x, y)$ and $F(z) = [A^T y; Ax]$, then

$$\langle F(z^*), z - z^* \rangle \geq 0 \quad \forall z \in \Delta_n \times \Delta_m =: C \quad (\text{VI})$$

Variational inequality

If $F = \nabla \varphi$ then (VI) would be equivalent to

$$\min_{z \in C} \varphi(z)$$

Potential - integrability

Question: Does there exist a potential φ for F , such that $F = \nabla\varphi$

Integrability condition (from calculus)

Is the case if

$$\frac{\partial\varphi}{\partial x\partial y} = \frac{\partial\varphi}{\partial y\partial x}$$

But

$$\frac{\partial F_1}{\partial y} = A^T \neq -A = \frac{\partial F_2}{\partial x}$$

VI as Fixed point equation

$$\begin{aligned}\langle F(z^*), z - z^* \rangle &\geq 0 \quad \forall z \in C \\ \Leftrightarrow z^* &= P_C(z^* - F(z^*))\end{aligned}\tag{FP}$$

Proof.

Applying the property of the projection

$$\langle P_C(x) - x, x' - P_C(x) \rangle \geq 0 \quad \forall x' \in C$$

with (FP), gives

$$\langle z^* - (z^* - F(z^*)), z - z^* \rangle \geq 0 \quad \forall z \in C. \quad \square$$

- ◇ should remind us of (projected) gradient descent
- ◇ when you see a fixed point equation: iterate!

But is it any good?

$$z_{k+1} = z_k - \alpha F(z_k)$$

Then

$$\begin{aligned}\|z_{k+1}\|^2 &= \|z_k\|^2 - \underbrace{2\alpha \langle F(z_k), z_k \rangle}_{=0} + \alpha^2 \|F(z_k)\|^2 \\ &= \|z_k\|^2 + \alpha^2 \|F(z_k)\|^2\end{aligned}$$

Resulting in $\|z_{k+1}\| \geq \|z_k\|$.

\Rightarrow **No bueno!**

We can still show

Theorem

complexity of $\mathcal{O}(1/\sqrt{k})$ for *averaged iterates* in terms of $f_p(x) - f_d(y)$.

- ◇ with the same analysis as for subgradient descent

Sketch of the proof

With the notation $g_k = F(z_k)$ we get

$$\begin{aligned}\|x_{k+1} - x^*\|^2 &\leq \|x_k - \alpha_k g_k - x^*\|^2 \\ &= \|x_k - x^*\|^2 + 2\alpha_k \langle g_k, x^* - x_k \rangle + \alpha^2 \|g_k\|^2.\end{aligned}$$

But now $\langle g_k, x^* - x_k \rangle = [f_d(y_k) - f_p(x_k)]$. Rest of the proof is left as an exercise.