

Mirror Descent

Axel Böhm

October 13, 2021

- 1 About norms
- 2 Bregman distances
- 3 Mirror descent

Recap on (sub)-gradient descent

- ◇ When we used a norm $\|\cdot\|$ we meant the 2-norm, i.e.

$$\|x\|_2 = \left(\sum_{i=1}^d x_i^2 \right)^{1/2}.$$

- ◇ In **gradient descent** we used Lipschitz continuity:

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|.$$

(Lead to a complexity of $\mathcal{O}(\frac{L}{\epsilon})$)

- ◇ For **sub-gradient descent** we used $\|g\| \leq G$ which lead to a complexity of $\mathcal{O}(\frac{G}{\epsilon})$.
- ◇ But there are **other norms**

$$\|x\|_\infty \leq \|x\|_2 \leq \sqrt{d}\|x\|_\infty$$

It can happen that $\|g\|_\infty \leq G$ but $\|g\|_2 \approx \sqrt{d}G$.

Recap on (sub)-gradient descent

- ◇ When we used a norm $\|\cdot\|$ we meant the 2-norm, i.e.

$$\|x\|_2 = \left(\sum_{i=1}^d x_i^2 \right)^{1/2}.$$

- ◇ In **gradient descent** we used Lipschitz continuity:

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|.$$

(Lead to a complexity of $\mathcal{O}(\frac{L}{\epsilon})$)

- ◇ For **sub-gradient descent** we used $\|g\| \leq G$ which lead to a complexity of $\mathcal{O}(\frac{G}{\sqrt{k}})$.
- ◇ But there are **other norms**

$$\|x\|_\infty \leq \|x\|_2 \leq \sqrt{d}\|x\|_\infty$$

It can happen that $\|g\|_\infty \leq G$ but $\|g\|_2 \approx \sqrt{d}G$.

Recap on (sub)-gradient descent

- ◇ When we used a norm $\|\cdot\|$ we meant the 2-norm, i.e.

$$\|x\|_2 = \left(\sum_{i=1}^d x_i^2 \right)^{1/2}.$$

- ◇ In **gradient descent** we used Lipschitz continuity:

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|.$$

(Lead to a complexity of $\mathcal{O}(\frac{L}{\epsilon})$)

- ◇ For **sub-gradient descent** we used $\|g\| \leq G$ which lead to a complexity of $\mathcal{O}(\frac{G}{\sqrt{k}})$.
- ◇ But there are **other norms**

$$\|x\|_\infty \leq \|x\|_2 \leq \sqrt{d}\|x\|_\infty$$

It can happen that $\|g\|_\infty \leq G$ but $\|g\|_2 \approx \sqrt{d}G$.

Different norms?

*But where did we use the norm in the **method**?*

Gradient Descent

$$x_{k+1} = x_k - \alpha_k \nabla f(x_k)$$

equivalently

$$x_{k+1} = \arg \min_{x \in \mathbb{R}^d} \left\{ f(x_k) + \langle \nabla f(x_k), x - x_k \rangle + \frac{1}{2\alpha_k} \|x - x_k\|_2^2 \right\}$$

We can replace the 2-norm with a more general **distance**.

Bregman distance

$h : \mathbb{R}^d \rightarrow \mathbb{R} \cup \{+\infty\}$ is convex

(i) h is differentiable on the interior of $\text{dom } h$

(ii) h is 1-strongly convex w.r.t. $\|\cdot\|_2$

Then

$$\mathcal{D}_h(x, y) := h(x) - h(y) - \langle \nabla h(y), x - y \rangle.$$

Properties

- ◇ $\mathcal{D}_h(x, y) \geq 0$
- ◇ $\mathcal{D}_h(x, y) \neq \mathcal{D}_h(y, x)$
- ◇ $\mathcal{D}_h(\cdot, y)$ is convex for all y

$$\mathcal{D}_h(x, y) \approx \frac{1}{2} \langle \nabla^2 h(y)(x - y), x - y \rangle = \frac{1}{2} \|x - y\|_{\nabla^2 h(y)}^2$$

- ◇ $\mathcal{D}_h(x, y) \geq \frac{1}{2} \|x - y\|^2$ (1-strong convexity)

Bregman distance

$h : \mathbb{R}^d \rightarrow \mathbb{R} \cup \{+\infty\}$ is convex

- (i) h is differentiable on the interior of $\text{dom } h$
- (ii) h is 1-strongly convex w.r.t. $\|\cdot\|_2$

Then

$$\mathcal{D}_h(x, y) := h(x) - h(y) - \langle \nabla h(y), x - y \rangle.$$

Properties

- ◇ $\mathcal{D}_h(x, y) \geq 0$
- ◇ $\mathcal{D}_h(x, y) \neq \mathcal{D}_h(y, x)$
- ◇ $\mathcal{D}_h(\cdot, y)$ is convex for all y

$$\mathcal{D}_h(x, y) \approx \frac{1}{2} \langle \nabla^2 h(y)(x - y), x - y \rangle = \frac{1}{2} \|x - y\|_{\nabla^2 h(y)}^2$$

- ◇ $\mathcal{D}_h(x, y) \geq \frac{1}{2} \|x - y\|^2$ (1-strong convexity)

Examples

- ◇ $h(x) = \frac{1}{2} \|x\|_2^2$ gives $\mathcal{D}_h(x, y) = \|x - y\|^2$
- ◇ $h(x) = \frac{1}{2(p-1)} \|x\|_p^2$ with $p \in [1, 2]$
- ◇ $\Delta^d = \{x \in \mathbb{R}_+^d : \sum_{i=1}^d x_i = 1\}$ the *unit simplex* and

$$h(x) = \begin{cases} \sum_{i=1}^d x_i \log(x_i) & x_i > 0 \\ +\infty & \text{otherwise} \end{cases}$$

the **Negative entropy**.

Negative entropy

- ◇ Negative entropy: $h(x) = \sum_{i=1}^d x_i \log(x_i)$ for $x_i > 0$.
- ◇ Then $\nabla h(x) = \log(x) + 1$ (coordinatewise) and

$$\begin{aligned}\mathcal{D}_h(x, y) &= \sum_{i=1}^d x_i \log(x_i) - y_i \log(y_i) - \langle \log(y) + 1, x - y \rangle \\ &= \sum_{i=1}^d x_i \log(x_i) - \sum_{i=1}^d x_i \log(y_i) \\ &= \sum_{i=1}^d x_i \log\left(\frac{x_i}{y_i}\right)\end{aligned}$$

Known as **Kullback-Leibler divergence** $K(X \| Y)$.

- ◇ Is strongly convex over Δ

$$\mathcal{D}(x, y) \geq \frac{1}{2} \|x - y\|_1^2 \quad \text{ Pinsker's ineq. }$$

Mirror descent

Idea: replace squared Euclidian norm with more general object:

$$\begin{aligned}x_{k+1} &= \arg \min_{x \in \mathbb{R}^d} \left\{ f(x_k) + \langle \nabla f(x_k), x - x_k \rangle + \frac{1}{\alpha_k} \mathcal{D}_h(x, x_k) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \nabla f(x_k), x \rangle + \frac{1}{\alpha_k} \mathcal{D}_h(x, x_k) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \nabla f(x_k), x \rangle + \frac{1}{\alpha_k} (h(x) - h(x_k) - \langle \nabla h(x_k), x - x_k \rangle) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \nabla f(x_k), x \rangle + \frac{1}{\alpha_k} (h(x) - \langle \nabla h(x_k), x \rangle) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \alpha_k \nabla f(x_k) - \nabla h(x_k), x \rangle + h(x) \right\}\end{aligned}$$

Question: But why *mirror* descent?

Mirror descent

Idea: replace squared Euclidian norm with more general object:

$$\begin{aligned}x_{k+1} &= \arg \min_{x \in \mathbb{R}^d} \left\{ f(x_k) + \langle \nabla f(x_k), x - x_k \rangle + \frac{1}{\alpha_k} \mathcal{D}_h(x, x_k) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \nabla f(x_k), x \rangle + \frac{1}{\alpha_k} \mathcal{D}_h(x, x_k) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \nabla f(x_k), x \rangle + \frac{1}{\alpha_k} (h(x) - h(x_k) - \langle \nabla h(x_k), x - x_k \rangle) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \nabla f(x_k), x \rangle + \frac{1}{\alpha_k} (h(x) - \langle \nabla h(x_k), x \rangle) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \alpha_k \nabla f(x_k) - \nabla h(x_k), x \rangle + h(x) \right\}\end{aligned}$$

Question: But why *mirror* descent?

Mirror descent

Idea: replace squared Euclidian norm with more general object:

$$\begin{aligned}x_{k+1} &= \arg \min_{x \in \mathbb{R}^d} \left\{ f(x_k) + \langle \nabla f(x_k), x - x_k \rangle + \frac{1}{\alpha_k} \mathcal{D}_h(x, x_k) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \nabla f(x_k), x \rangle + \frac{1}{\alpha_k} \mathcal{D}_h(x, x_k) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \nabla f(x_k), x \rangle + \frac{1}{\alpha_k} (h(x) - h(x_k) - \langle \nabla h(x_k), x - x_k \rangle) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \nabla f(x_k), x \rangle + \frac{1}{\alpha_k} (h(x) - \langle \nabla h(x_k), x \rangle) \right\} \\&= \arg \min_{x \in \mathbb{R}^d} \left\{ \langle \alpha_k \nabla f(x_k) - \nabla h(x_k), x \rangle + h(x) \right\}\end{aligned}$$

Question: But why *mirror* descent?

The Mirror part

$$x_{k+1} = \arg \min_{x \in \mathbb{R}^d} \{ \langle \alpha_k \nabla f(x_k) - \nabla h(x_k), x \rangle + h(x) \}$$

By optimality condition:

$$0 = \alpha_k \nabla f(x_k) - \nabla h(x_k) + \nabla h(x_{k+1})$$

Therefore

$$\nabla h(x_{k+1}) = \nabla h(x_k) - \alpha_k \nabla f(x_k)$$

Why it's called mirror descent




Figure: $\psi = h$

Mirror Descent on the unit simplex

Negative entropy: $h(x) = \sum_{i=1}^d x_i \log(x_i)$ for $x_i > 0$.

We define $a := \alpha_k \nabla f(x_k) - \nabla h(x_k)$. Then

$$x_{k+1} = \arg \min_{x \in \Delta} \{ \langle a, x \rangle + h(x) \}$$

with $x_i \geq 0$ and $\sum x_i = 1$.

How to solve this?

Via **Lagrange**

$$L(x, \mu) = \langle a, x \rangle + h(x) - \mu(x_1 + \cdots + x_d - 1)$$

Mirror Descent on the unit simplex [contd]

Then,

$$\partial_{x_i} L(x, \mu) = a_i + \log(x_i) + 1 - \mu \stackrel{!}{=} 0$$

$$\log(x_i) = \mu - 1 - a_i$$

$$x_i = e^{\mu-1-a_i} = \beta e^{-a_i}$$

with $\beta = e^{\mu-1}$.

Second constraint:

$$\sum_{i=1}^d x_i \stackrel{!}{=} 1 \Rightarrow \sum_{i=1}^d \beta e^{-a_i} = 1 \Rightarrow \beta = \frac{1}{\sum_{i=1}^d e^{-a_i}} \Rightarrow x_i = \frac{e^{-a_i}}{\sum_{j=1}^d e^{-a_j}}$$

Final mirror descent update:

$$x_{k+1}(i) = \frac{x_k(i) e^{\alpha_k [\nabla f(x_k)]_i}}{\sum_{j=1}^d e^{\alpha_k [\nabla f(x_k)]_j}}$$

Mirror Descent on the unit simplex [contd]

Then,

$$\partial_{x_i} L(x, \mu) = a_i + \log(x_i) + 1 - \mu \stackrel{!}{=} 0$$

$$\log(x_i) = \mu - 1 - a_i$$

$$x_i = e^{\mu-1-a_i} = \beta e^{-a_i}$$

with $\beta = e^{\mu-1}$.

Second constraint:

$$\sum_{i=1}^d x_i \stackrel{!}{=} 1 \Rightarrow \sum_{i=1}^d \beta e^{-a_i} = 1 \Rightarrow \beta = \frac{1}{\sum_{i=1}^d e^{-a_i}} \Rightarrow x_i = \frac{e^{-a_i}}{\sum_{j=1}^d e^{-a_j}}$$

Final mirror descent update:

$$x_{k+1}(i) = \frac{x_k(i) e^{\alpha_k [\nabla f(x_k)]_i}}{\sum_{j=1}^d e^{\alpha_k [\nabla f(x_k)]_j}}$$

Mirror Descent on the unit simplex [contd]

Then,

$$\partial_{x_i} L(x, \mu) = a_i + \log(x_i) + 1 - \mu \stackrel{!}{=} 0$$

$$\log(x_i) = \mu - 1 - a_i$$

$$x_i = e^{\mu-1-a_i} = \beta e^{-a_i}$$

with $\beta = e^{\mu-1}$.

Second constraint:

$$\sum_{i=1}^d x_i \stackrel{!}{=} 1 \Rightarrow \sum_{i=1}^d \beta e^{-a_i} = 1 \Rightarrow \beta = \frac{1}{\sum_{i=1}^d e^{-a_i}} \Rightarrow x_i = \frac{e^{-a_i}}{\sum_{j=1}^d e^{-a_j}}$$

Final mirror descent update:

$$x_{k+1}(i) = \frac{x_k(i) e^{\alpha_k [\nabla f(x_k)]_i}}{\sum_{j=1}^d e^{\alpha_k [\nabla f(x_k)]_j}}$$

(General) mirror descent convergence statement

$$\|y\|_* := \max_{\|x\|=1} \{\langle y, x \rangle\}$$

Theorem

In $(\mathbb{R}^d, \|\cdot\|)$ and subgradients bounded in dual norm $\|g_k\|_ \leq G$, then*

$$f(\hat{x}_k) - f^* \leq \frac{\mathcal{D}(x^*, x_0) + \frac{1}{2} \sum_i \alpha_i^2 G^2}{\sum_i \alpha_i}$$

Convergence on the unit simplex

What about $\mathcal{D}(x^*, x_0)$? Let $x_0 = (\frac{1}{n}, \dots, \frac{1}{n})$, then

$$\mathcal{D}(x, x_0) = \sum x_i \log \left(\frac{x_i}{\frac{1}{n}} \right) = \sum x_i \log(x_i) + \log(n) \leq \log(n)$$

while $\|x_0 - x^*\|^2 \leq 2$.

But if

$$\|g\|_\infty = \|g\|_1^* \leq G$$

we can still have

$$\|g\|_2 \approx \sqrt{d}G.$$