

Theorem[section] Corollary[section] Lemma[section]
Definition[section]

Online Optimization

Axel

September 3, 2021

Introduction

What is online Learning

Consider the following repeated game:

In each round $t = 1, \dots, T$

- ▶ An adversary choose a real number in $y_t \in [0, 1]$ and he keeps it secret;
- ▶ You try to guess the real number, choosing $x_t \in [0, 1]$;
- ▶ The adversary's number is revealed and you pay the squared difference $(x_t - y_t)^2$.

Task: guess a sequence of numbers as precisely as possible. To be a game, we now have to decide what is the “winning condition”. Let's see what makes sense to consider as winning condition.

Question: How to measure success?

Adversary plays i.i.d.

Consider: Adversary number are drawn from a fixed distribution (with mean μ and Variance σ^2). If we knew the distribution, we could pick the mean and pay in expectation $\sigma^2 T$ (optimal).

$$\mathbb{E}_Y \left[\sum_{t=1}^T (x_t - Y)^2 \right] - \sigma^2 T,$$

or equivalently considering the average

$$\frac{1}{T} \mathbb{E}_Y \left[\sum_{t=1}^T (x_t - Y)^2 \right] - \sigma^2 .$$

Minimizing Regret

Let's rewrite a bit more general

$$\mathbb{E} \left[\sum_{t=1}^T (x_t - Y)^2 \right] - \min_{x \in [0,1]} \mathbb{E} \left[\sum_{t=1}^T (x - Y)^2 \right] .$$

($\sigma^2 T$ was nothing other than the payoff of the best possible strategy)

Finally: remove the assumption on how the data is generated, consider any arbitrary sequence of y_t (we can remove the expectation because there is no stochasticity anymore).

$$R := \sum_{t=1}^T (x_t - y_t)^2 - \min_{x \in [0,1]} \sum_{t=1}^T (x - y_t)^2$$

The quantity above is called the *Regret*, because it measures how much the algorithm regrets for not sticking on all the rounds to the optimal choice in hindsight.

General loss functions

Online Learning is nothing else than designing and analyzing algorithms to minimize the Regret over a sequence of loss functions with respect to an arbitrary competitor $\in V \subseteq \mathbb{R}^d$:

$$\tau() := \sum_{t=1}^T \ell_t(t) - \sum_{t=1}^T \ell_t() .$$

This framework is pretty powerful, and it allows to reformulate a bunch of different problems in machine learning and optimization as similar games. More in general, with the regret framework we can analyze situations in which the data are not independent and identically distributed from a distribution, yet I would like to guarantee that the algorithm is “learning” something. For example, online learning can be used to analyze

- ▶ Click prediction problems;
- ▶ Routing on a network;
- ▶ Convergence to equilibrium of repeated games.

It can *also* be used to analyze stochastic algorithms, e.g., Stochastic Gradient Descent, but the adversarial nature of the

Back to the numbers game

Let's take a look at the best strategy in hindsight, that is argmin of the second term of the regret. It should be immediate to see that

$$x_T^* := \arg \min_{x \in [0,1]} \sum_{t=1}^T (x - y_t)^2 = \frac{1}{T} \sum_{t=1}^T y_t .$$

Since we do not know the future x_T^* is not an option in each round. However, we do know the past, so a reasonable strategy in each round could be to output the best number over the past.

For sure, the reason why it could work is not because we expect the future to be like the past, because it is not true! Instead, we want to leverage the fact that the optimal guess should not change too much between rounds, so we can try to “track” it over time.

Hence, on each round t our strategy is to guess

$x_t = x_{t-1}^* = \frac{1}{t-1} \sum_{i=1}^{t-1} y_i$. Such strategy is usually called *Follow-the-Leader* (FTL), because you are following what would have been the optimal thing to do on the past rounds (the Leader).

Follow the leader

Let's now try to show that indeed this strategy will allow us to win the game.

Lemma

Let $V \subseteq \mathbb{R}^d$ and $\ell_t : V \rightarrow \mathbb{R}$ an arbitrary sequence of loss functions. Denote by x_t^ a minimizer of the cumulative losses over the previous t rounds in V . Then, we have*

$$\sum_{t=1}^T \ell_t(x_t^*) \leq \sum_{t=1}^T \ell_t(x_T^*) .$$

Proof.

We prove it by induction over T . The base case is

$$\ell_1(x_1^*) \leq \ell_1(x_1^*),$$

that is trivially true. Now, for $T \geq 2$, we assume that $\sum_{t=1}^{T-1} \ell_t(x_t^*) \leq \sum_{t=1}^{T-1} \ell_t(x_{T-1}^*)$ is true and we must prove the stated inequality, that is

Follow the leader II

Theorem

Let $y_t \in [0, 1]$ for $t = 1, \dots, T$ an arbitrary sequence of numbers. Let the algorithm's output $x_t = x_{t-1}^* := \frac{1}{t-1} \sum_{i=1}^{t-1} y_i$. Then, we have

$$T = \sum_{t=1}^T (x_t - y_t)^2 - \min_{x \in [0,1]} \sum_{t=1}^T (x - y_t)^2 \leq 4 + 4 \ln T .$$

Proof.

Exercise. □

Failure of FTL

Let $V = [-1, 1]$ and consider the sequence of losses

$\ell_t(x) = z_t x + i_V(x)$, where $z_1 = -0.5$,

$z_t = 1, t = 2, 4, \dots$

$z_t = -1, t = 3, 5, \dots$ Predictions of FTL will be $x_t = 1$ for t even and $x_t = -1$ for t odd. Cumulative loss of the FTL algorithm will be T while the cumulative loss of the fixed solution $u = 0$ is 0.

Thus, the regret of FTL is T .

Outlook: Online gradient descent will be optimal

Weighted majority algorithm

Consider the *learning from experts* scenario. Experts = $1, \dots, n$.
Decision: “Yes” or “No”.

$$f_t(x_t) = \begin{cases} 1 & \text{if wrong} \\ 0 & \text{otherwise} \end{cases}$$

1. $w_1(i) = 1$ for all $i = 1, \dots, n$
2. for $t = 1, \dots, T$
 - 2.1 compare weights $\sum_{i \in YES} w_t(i)$ vs. $\sum_{i \in NO} w_t(i)$
 - 2.2 choose Yes or No depending on above comparison
 - 2.3 observe feedback
 - 2.4 update weights:

$$w_{t+1}(i) = \begin{cases} w_t(i) & \text{if Expert } i \text{ was right} \\ (1 - \alpha)w_t(i) & \text{if Expert } i \text{ made a mistake} \end{cases}$$

Weighted majority algorithm II

Theorem

Let M_t be the number mistakes we make after t attempts and $m_t(i) = \#$ the number of mistakes expert i made... Then,

$$M_T \leq 2(1 + \alpha)m_T(i) + 2\frac{\log(n)}{\alpha}$$

Also

$$M_T - m_T(i^*) = R_T$$

Proof of the Theorem

We always have $\|w_{t+1}\|_1 \leq \|w_t\|_1$. Also, if we made a mistake, then

$$\begin{aligned}\|w_{t+1}\|_1 &\leq \frac{1}{2}\|w_t\|_1 + \frac{1}{2}\|1w_t\|(1-\alpha) \\ &= \|w_t\|_1(1-\alpha/2) \\ &\leq \|w_1\|_1(1-\alpha/2)^{M_t} \\ &= n(1-\alpha/2)^{M_t}\end{aligned}$$

Next

$$w_{t+1}(i) = (1-\alpha)^{m_t(i)} \leq \|w_{t+1}\|_1$$

Combining the above two yields

$$(1-\alpha)^{m_t(i)} \leq n(1-\alpha/2)^{M_t}$$

and

$$m_t(i) \log(1-\alpha) \leq \log(n) + M_T \log(1-\alpha/2).$$

remainder of the proof

use the fact that for $x \in (0, \frac{1}{2})$

$$-x - x^2 \leq \log(1 - x) \leq -x$$

to deduce that

$$-m_t(i)(\alpha + \alpha^2) \leq \log(n) - M_T \frac{\alpha}{2} - 2m_t(i)(1 + \alpha) \leq \frac{2}{\alpha} \log(n) - M_T$$

which yields

$$M_T - \leq 2m_t(i)(1 + \alpha) + \frac{2}{\alpha} \log(n) -$$

Randomized Weighted Majority

Instead of picking the opinion of the (weighted) majority, we only do so with a **probability**.

1. $w_1(i) = 1$ for all $i = 1, \dots, n$ and $\alpha \in (0, 1)$
2. for $t = 1, \dots, T$
 - 2.1 compute $p_t(i) = w_t(i) / \|w_t\|_1$
 - 2.2 choose expert i with probability $p_t(i)$
 - 2.3 observe feedback
 - 2.4 update weights:

$$w_{t+1}(i) = \begin{cases} w_t(i) & \text{if expert } i \text{ was right} \\ (1 - \alpha)w_t(i) & \text{if expert } i \text{ made a mistake} \end{cases}$$

Randomized Weighted Majority

Instead of picking the option of the (weighted) majority, we only do so with a **probability**.

1. $w_1(i) = 1$ for all $i = 1, \dots, n$ and $\alpha \in (0, 1)$
2. for $t = 1, \dots, T$
 - 2.1 compute $p_t(i) = w_t(i) / \|w_t\|_1$
 - 2.2 choose expert i with probability $p_t(i)$
 - 2.3 observe feedback
 - 2.4 update weights:

$$w_{t+1}(i) = \begin{cases} w_t(i) & \text{if expert } i \text{ was right} \\ (1 - \alpha)w_t(i) & \text{if expert } i \text{ made a mistake} \end{cases}$$

Comment: Randomizing algorithms typically improves the (worst case) analysis.

Randomized Weighted Majority

As before:

$M_t = \#$ of mistakes we make after t attempts and $m_t(i) = \#$ of mistakes expert i made.

Theorem

$$\mathbb{E}[M_T] \leq (1 + \alpha)m_T(i) + \frac{\log(n)}{\alpha}$$

Where did we improve? somewhere in the constants.