

Online Optimization

Axel

September 7, 2021

1 Introduction

2 Strategies

What is online Learning

Consider the following repeated game:

In each round $t = 1, \dots, T$

- An adversary choose a real number in $y_t \in [0, 1]$ and he keeps it secret;
- You try to guess the real number, choosing $x_t \in [0, 1]$;
- The adversary's number is revealed and you pay the squared difference $(x_t - y_t)^2$.

Task: guess a sequence of numbers as precisely as possible. To be a game, we now have to decide what is the “winning condition”.

Let's see what makes sense to consider as winning condition.

Question: How to measure success?

Adversary plays i.i.d.

Consider: Adversary number are drawn from a fixed distribution (with mean μ and Variance σ^2). If we knew the distribution, we could pick the mean and pay in expectation $\sigma^2 T$ (optimal).

$$\mathbb{E}_Y \left[\sum_{t=1}^T (x_t - Y)^2 \right] - \sigma^2 T,$$

or equivalently considering the average

$$\frac{1}{T} \mathbb{E}_Y \left[\sum_{t=1}^T (x_t - Y)^2 \right] - \sigma^2 .$$

Minimizing Regret

Let's rewrite a bit more general

$$\mathbb{E} \left[\sum_{t=1}^T (x_t - Y)^2 \right] - \min_{x \in [0,1]} \mathbb{E} \left[\sum_{t=1}^T (x - Y)^2 \right] .$$

($\sigma^2 T$ was nothing other than the payoff of the best possible strategy)

Finally: remove the assumption on how the data is generated, consider any arbitrary sequence of y_t (we can remove the expectation because there is no stochasticity anymore).

$$R_T := \sum_{t=1}^T (x_t - y_t)^2 - \min_{x \in [0,1]} \sum_{t=1}^T (x - y_t)^2$$

The quantity above is called the *Regret*, because it measures how much the algorithm regrets for not sticking on all the rounds to the

General loss functions

Online Learning is the study of algorithms to minimize the *regret* over a sequence of loss functions with respect to an arbitrary competitor $u \in V \subseteq \mathbb{R}^d$:

$$R_T(u) := \sum_{t=1}^T \ell_t(x_t) - \sum_{t=1}^T \ell_t(u) .$$

Regret framework allows to

- reformulate problems in machine learning and optimization as similar games.
- analyze situations in which the data are not i.i.d. yet want to guarantee that the algorithm is “learning” something.

For example, online learning can be used to analyze

- Click prediction problems;
- Routing on a network;

Back to the numbers game

Let's take a look at the **best strategy in hindsight**, that is argmin of the second term of the regret. Clearly

$$x_T^* := \arg \min_{x \in [0,1]} \sum_{t=1}^T (x - y_t)^2 = \frac{1}{T} \sum_{t=1}^T y_t .$$

- Don't know the future: x_T^* is not an option in each round
- But do know the past. in each round: best number over the past.
- not because we expect the future to be like the past (not true)
- optimal guess should not change too much between rounds (so we can try to “track” it over time)

Hence, on each round t our strategy is to guess

$x_t = x_{t-1}^* = \frac{1}{t-1} \sum_{i=1}^{t-1} y_i$, called **Follow-the-Leader (FTL)**.

Follow the leader

Let's now try to show that indeed this strategy will allow us to win the game.

Lemma

Let $V \subseteq \mathbb{R}^d$ and $\ell_t : V \rightarrow \mathbb{R}$ an arbitrary sequence of loss functions. Denote by x_t^ a minimizer of the cumulative losses over the previous t rounds in V . Then, we have*

$$\sum_{t=1}^T \ell_t(x_t^*) \leq \sum_{t=1}^T \ell_t(x_T^*) .$$

Proof.

We prove it by induction over T . The base case is



Follow the leader II

Theorem

Let $y_t \in [0, 1]$ for $t = 1, \dots, T$ an arbitrary sequence of numbers.
Let the algorithm's output $x_t = x_{t-1}^* := \frac{1}{t-1} \sum_{i=1}^{t-1} y_i$. Then, we have

$$R_T = \sum_{t=1}^T (x_t - y_t)^2 - \min_{x \in [0,1]} \sum_{t=1}^T (x - y_t)^2 \leq 4 + 4 \ln T .$$

Proof.

Exercise. □

Failure of FTL

Let $V = [-1, 1]$ and consider the sequence of losses

$$\ell_t(x) = z_t x + i_V(x), \text{ where } z_1 = -0.5 \quad z_t = \begin{cases} 1, & t \text{ even} \\ -1, & t \text{ odd} \end{cases}$$

Predictions of FTL will be $x_t = 1$ for t even and $x_t = -1$ for t odd. Cumulative loss of the FTL algorithm will be T while the cumulative loss of the fixed solution $u = 0$ is 0. Thus, the regret of FTL is T .

Outlook:

- Follow the *regularized* leader
- Online gradient descent

Weighted majority algorithm

Consider the *learning from experts* scenario. Experts = $1, \dots, n$.
Decision: “Yes” or “No”.

$$f_t(x_t) = \begin{cases} 1 & \text{if wrong} \\ 0 & \text{otherwise} \end{cases}$$

- 1 $w_1(i) = 1$ for all $i = 1, \dots, n$
- 2 for $t = 1, \dots, T$
 - 1 compare weights $\sum_{i \in YES} w_t(i)$ vs. $\sum_{i \in NO} w_t(i)$
 - 2 choose Yes or No depending on above comparison
 - 3 observe feedback
 - 4 update weights:

$$w_{t+1}(i) = \begin{cases} w_t(i) & \text{if Expert } i \text{ was right} \\ (1 - \alpha)w_t(i) & \text{if Expert } i \text{ made a mistake} \end{cases}$$

Weighted majority algorithm II

Theorem

Let M_t be the number mistakes we make after t attempts and $m_t(i) = \#$ the number of mistakes expert i made... Then,

$$M_T \leq 2(1 + \alpha)m_T(i) + 2\frac{\log(n)}{\alpha}$$

Also

$$M_T - m_T(i^*) = R_T$$

Proof of the Theorem

We always have $\|w_{t+1}\|_1 \leq \|w_t\|_1$. Also, if we made a mistake, then

$$\begin{aligned}\|w_{t+1}\|_1 &\leq \frac{1}{2}\|w_t\|_1 + \frac{1}{2}\|1w_t\|(1-\alpha) \\ &= \|w_t\|_1(1-\alpha/2) \\ &\leq \|w_1\|_1(1-\alpha/2)^{M_t} \\ &= n(1-\alpha/2)^{M_t}\end{aligned}$$

Next

$$w_{t+1}(i) = (1-\alpha)^{m_t(i)} \leq \|w_{t+1}\|_1$$

Combining the above two yields

$$(1-\alpha)^{m_t(i)} \leq n(1-\alpha/2)^{M_t}$$

and

$$m_t(i) \log(1-\alpha) \leq \log(n) + M_t \log(1-\alpha/2)$$

remainder of the proof

use the fact that for $x \in (0, \frac{1}{2})$

$$-x - x^2 \leq \log(1 - x) \leq -x$$

to deduce that

$$-m_t(i)(\alpha + \alpha^2) \leq \log(n) - M_T \frac{\alpha}{2} - 2m_t(i)(1 + \alpha) \leq \frac{2}{\alpha} \log(n) - M_T$$

which yields

$$M_T - \leq 2m_t(i)(1 + \alpha) + \frac{2}{\alpha} \log(n) -$$

Randomized Weighted Majority

Instead of picking the opinion of the (weighted) majority, we only do so with a **probability**.

- 1 $w_1(i) = 1$ for all $i = 1, \dots, n$ and $\alpha \in (0, 1)$
- 2 for $t = 1, \dots, T$
 - 1 compute $p_t(i) = w_t(i) / \|w_t\|_1$
 - 2 choose expert i with probability $p_t(i)$
 - 3 observe feedback
 - 4 update weights:

$$w_{t+1}(i) = \begin{cases} w_t(i) & \text{if expert } i \text{ was right} \\ (1 - \alpha)w_t(i) & \text{if expert } i \text{ made a mistake} \end{cases}$$

Randomized Weighted Majority

Instead of picking the opinion of the (weighted) majority, we only do so with a **probability**.

- 1 $w_1(i) = 1$ for all $i = 1, \dots, n$ and $\alpha \in (0, 1)$
- 2 for $t = 1, \dots, T$
 - 1 compute $p_t(i) = w_t(i) / \|w_t\|_1$
 - 2 choose expert i with probability $p_t(i)$
 - 3 observe feedback
 - 4 update weights:

$$w_{t+1}(i) = \begin{cases} w_t(i) & \text{if expert } i \text{ was right} \\ (1 - \alpha)w_t(i) & \text{if expert } i \text{ made a mistake} \end{cases}$$

Comment: Randomizing algorithms typically improves the (worst case) analysis.

Randomized Weighted Majority contd.

As before:

M_t = # of mistakes we make after t attempts and $m_t(i)$ = # of mistakes expert i made.

Theorem

$$\mathbb{E}[M_T] \leq (1 + \alpha) m_T(i) + \frac{\log(n)}{\alpha}$$

Improved constants!

proof of randomized WMA

Multiplicative Weights Algorithm

Before: Loss l_t was 0 or 1 Now: General loss functions

$\ell_t = (\ell_t(1), \dots, \ell_t(n))$ with $\ell_t(i) \in [-1, 1]$

1 $w_1(i) = 1$ for all $i = 1, \dots, n$ and $\alpha \in (0, 1)$

2 for $t = 1, \dots, T$

1 compute $p_t(i) = w_t(i) / \|w_t\|_1$

2 choose expert i with probability $p_t(i)$

3 observe loss $\ell_t = (\ell_t(1), \dots, \ell_t(n))$

4 update weights:

$$w_{t+1}(i) = (1 - \alpha \ell_t(i)) w_t(i)$$

Note that

$$\langle \mathbf{p}_t, \ell_t \rangle = p_t(1)\ell_t(1) + \dots + p_t(n)\ell_t(n) = \mathbb{E}_i[\ell_t(i)]$$

gives expected loss of round t .

Multiplicative Weights Algorithm [contd]

Theorem

if $\ell_t(i) \in [-1, 1]$ and $\alpha < \frac{1}{2}$, then **MWA** guarantees

$$\sum_{t=1}^T \langle \mathbf{p}_t, \ell_t \rangle - \sum_{t=1}^T \ell_t(i) \leq \alpha \sum_{t=1}^T |\ell_t(i)| + \frac{\log(n)}{\alpha} \quad \forall i$$

Hedge Algorithm

- 1 $w_1(i) = 1$ for all $i = 1, \dots, n$ and $\alpha \in (0, 1)$
- 2 for $t = 1, \dots, T$
 - 1 compute $p_t(i) = w_t(i) / \|w_t\|_1$
 - 2 choose expert i with probability $p_t(i)$
 - 3 observe loss $\ell_t = (\ell_t(1), \dots, \ell_t(n))$
 - 4 update weights:

$$w_{t+1}(i) = w_t(i) e^{-\alpha \ell_t(i)}$$

Note:

$$e^{-x} \approx 1 - x$$

Hedge Algorithm [contd]

Theorem

If $\ell_t(i) \in [-1, 1]$ and $\alpha < \frac{1}{2}$, then **Hedge Alg.** guarantees

$$\sum_{t=1}^T \langle \mathbf{p}_t, \ell_t \rangle - \sum_{t=1}^T \ell_t(i) \leq \alpha \sum_{t=1}^T \langle \mathbf{p}_t, \ell_t \rangle^2 + \frac{\log(n)}{\alpha} \quad \forall i.$$

Hedge Algorithm [contd]

Theorem

If $\ell_t(i) \in [-1, 1]$ and $\alpha < \frac{1}{2}$, then **Hedge Alg.** guarantees

$$\sum_{t=1}^T \langle \mathbf{p}_t, \ell_t \rangle - \sum_{t=1}^T \ell_t(i) \leq \alpha \sum_{t=1}^T \langle \mathbf{p}_t, \ell_t \rangle^2 + \frac{\log(n)}{\alpha} \quad \forall i.$$

Observe: Iteration t is just

$$w_{t+1}(i) = w_t(i) e^{-\alpha \ell_t(i)}$$

$$p_{t+1}(i) = \frac{w_{t+1}(i)}{\|\mathbf{w}_{t+1}\|_1}$$

Online mirror descent! (KL-divergence setting:)

$$h(\mathbf{x}) = \sum_i x(i) \log(x(i))$$