

Variance reduction for stochastic gradient methods

Axel Böhm

September 13, 2021

1 Introduction

2 SAG

3 SAGA

4 SVRG

A common Task in (supervised) machine learning:

$$\min_{x \in \mathbb{R}^d} f(x) := \frac{1}{n} \sum_{i=1}^n \underbrace{\text{loss for } i\text{-th sample}} + \underbrace{\psi(x)}_{\text{regularizer}}$$

where the i -th sample is (a_i, y_i) .

- linear regression: $f_i(x) = (a_i^T x - y_i)^2$, and $\psi = 0$
- logistic regression: $f_i(x) = \log(1 + e^{-y_i a_i^T x})$, and $\psi = 0$
“sigmoid function” and logistic loss.
- Lasso: f_i as for linear regression but $\psi(x) = \|x\|_1$
- SVM: $f_i(x) = \max\{0, 1 - y_i a_i^T x\}$ and $\psi(x) = \|x\|^2$

Stochastic gradient descent

Algorithm SGD

- 1: **for** $k = 1, 2, \dots$ **do**
 - 2: pick i_k uniform at random in $[n]$
 - 3: $x_{k+1} = x_k - \alpha_k \nabla f_{i_k}(x_k)$
-

We already noticed that:

- large stepsizes fail to suppress variability in the stoch. gradients leads to oscillations
- decreasing stepsizes mitigate this problem but slows down convergence (too conservative)

Recall SGD (template)

$$x_{k+1} = x_k - \alpha_k g_k$$

- g_k is an unbiased estimator of the true gradient $\nabla F(x_k)$
- convergence depends on variance $\mathbb{E}[\|g_k - \nabla F(x_k)\|] \leq \sigma_g$
- vanilla SGD $g_k = \nabla f_{i_k}(x_k)$
issue: σ_g is non-negligible even close to the solution
- **Q:** can we choose g_k in a different way to reduce variability?

A simple idea

Consider

- estimator X for parameter μ ($\mathbb{E}[X] = \mu$ and $\mathbb{V}[X] = \sigma^2$)
- want to keep unbiased but reduce variance
- find Y such that $\mathbb{E}[Y] = 0$ but $\text{Cov}(X, Y)$ is large and $\tilde{X} = X - Y$
- remains unbiased

$$\mathbb{V}[\tilde{X}] = \mathbb{V}[X] + \mathbb{V}[Y] - 2 \text{Cov}[X, Y]$$

- can be much smaller than $\mathbb{V}[X]$ if X, Y are highly correlated

Stochastic average gradient (SAG), 2013

- maintain table containing gradients g_i of f_i
- at step $k = 1, 2, \dots$ pick random $i_k \in [n]$ and

$$g_{i_k}^k := \nabla f_{i_k}(x^k)$$

set $g_i^k = g_i^{k-1}$ for all $i \neq i_k$ (remain the same)

- Update

$$x^{k+1} = x^k - \alpha_k \frac{1}{n} \sum_{i=1}^n g_i^k$$

- assuming gradients do not change too much along trajectory
- gradient estimator **no longer unbiased**
- Isn't it expensive to average these gradients?

$$x^{k+1} = x^k - u\alpha_k \left(\frac{g_{i_k}^k}{n} - \frac{g_{i_k}^{k-1}}{n} + \underbrace{\frac{1}{n} \sum_{i=1}^n g_i^k}_{\text{old table average}} \right)$$

SAG variance reduction

Gradient estimator in SAG:

$$x^{k+1} = x^k - \alpha_k \frac{1}{n} \left(\underbrace{g_{i_k}^k}_X - \underbrace{g_{i_k}^{k-1} - \sum_{i=1}^n g_i^k}_Y \right)$$

- While $\mathbb{E}[X] = \nabla f(x^k)$, but $\mathbb{E}[Y] \neq 0 \rightarrow$ biased estimator
- X and Y are correlated as $X - Y \rightarrow 0$ as
- x^k and x_{k-1} both converge to x^* and
- the last term converges to $\nabla f(x^*) = 0$

Convergence

As always, initialization plays a role: $D^2 := \|x^0 - x^*\|^2$.


$$\text{SAG: } \frac{n}{k}(f(x^0) - f^*) + \frac{L}{k}D^2$$

$$\text{GD: } \frac{L}{k}D^2$$

$$\text{SAG: } \frac{L}{\sqrt{k}}D^2$$

- Achieves linear convergence in the *strongly* convex setting.
- proofs are difficult (and computer-aided)

Same gradient oracle cost as SGD, but same converge rate as GD.



SAGA, 2014

Very similar:

- maintain table containing gradients g_i of f_i
- at step $k = 1, 2, \dots$ pick random $i_k \in [n]$ and

$$g_{i_k}^k := \nabla f_{i_k}(x^k)$$

set $g_i^k = g_i^{k-1}$ for all $i \neq i_k$ (remain the same)

- Update

$$x^{k+1} = x^k - u\alpha_k \left(g_{i_k}^k - g_{i_k}^{k-1} + \frac{1}{n} \sum_{i=1}^n g_i^k \right)$$

- estimator now **unbiased!**


For Comparison

SAGA gradient estimate:

$$g_{i_k}^k - g_{i_k}^{k-1} + \frac{1}{n} \sum_{i=1}^n g_i^k$$

SAGA gradient estimate:

$$\frac{1}{n} g_{i_k}^k - \frac{1}{n} g_{i_k}^{k-1} + \frac{1}{n} \sum_{i=1}^n g_i^k$$



Stochastic Variance Reduced Gradient (SVRG), 2013

Algorithm SVRG

```
1: for  $k = 1, 2, \dots$  do
2:   Set  $x^0 = \tilde{x} = \tilde{x}^k$ 
3:   Compute  $\tilde{\mu} := \nabla f(\tilde{x})$ 
4:   for  $l = 1, 2, \dots, m$  do
5:     pick  $i_l$  uniform at random in  $[n]$ 
6:     Set  $x^{l+1} = x^l - \alpha(\nabla f_{i_l}(x^l) - \nabla f_{i_l}(\tilde{x}) + \tilde{\mu})$ 
```

- Does not need to store full table of gradients.
- requires *batch* gradient computation every *epoch*
- convergence rates similar to SAGA, but simpler analysis.
- per iteration cost of SVRG is comparable to that of SGD if $m \geq n$